

Computer Vision I

Assignment 1

Prof. Stefan Roth
Krishnakant Singh
Shweta Mahajan



TECHNISCHE
UNIVERSITÄT
DARMSTADT

08/11/2021

This assignment is due on November 22nd, 2021 at 23:59.

Group work and grading policy

You are required to work on each assignment in groups of *three*. We reserve the right to merge singleton groups randomly for subsequent assignments.

Programming exercises

For the programming exercises you will be asked to hand in Python code. Please use the provided environment as we will use this environment to test your solution. Use comments in the source code to clarify the functionality of your implementation in sufficient detail. Even if the computed results are incorrect due to a minor bug, you may still earn partial credit if your reasoning is valid.

You *must* adhere to the naming scheme for functions and files included with each problem. Do *not* alter function files and do *not* change the given function signatures. If you feel that there is a mistake in the assignments, or you find the task ambiguous, contact us on Moodle.

Multiple choice questions

There may be multiple choice questions in the assignments, which are implemented via a function or a class method. We provide a detailed explanation *in the source code* about the data type and format of the return value. Please, read the instructions carefully, because the format may vary depending on the question.

Files you need

The data and source code skeleton and the PDF with the assignment tasks will be made available on Moodle.

Handing in

Please, upload your solutions in the corresponding section on Moodle. Each problem task will specify the files to be included in your submission. Please include only the files specified in the problem descriptions and do *not* upload data – such as images – provided by us. You only need to submit one solution per group. Should you have troubles accessing Moodle, get in touch with us as soon as possible. Upload all your solution files as a single `.zip` or `.tar.gz` file. *We do not accept other file formats!*

Late submissions

We will accept late submissions, but you will lose *20% of the total reachable points* for every day of delay. Note that even 15 minute delays will be counted as being one day late! After the assignment solution has been discussed in class, you may no longer hand in.

Other remarks

Your grade will depend on two factors. Firstly, it will be determined by the correctness of your answer. Secondly, it will depend on the clarity of presenting your results and a good writing style. It is your task to find a way to *explain clearly how* you solved the problems. You can still get partial credit even if you did not complete the task.

We encourage interaction about class-related topics both in-class and on Moodle. However, you should not share solutions with other groups, and **everything you hand in must be your own work**. You are also not allowed to copy material from the web without acknowledgment. You must **acknowledge any source of information that you used to solve the homework** (i.e. books other than the course books, papers, etc.). Using acknowledgments will *not* affect your grade, but failing to do so is a clear violation of academic ethics. Note that the university as well as the department is very serious about plagiarism. For more details please see <https://plagiarism.org>.

Problem 1: Getting to know Python (5 Points)

In this task you will set up Python-3.8 on your system using Miniconda and learn how to install additional packages. Packages that you will often use are: (i) Pillow for loading and saving images; (ii) NumPy and SciPy for scientific computing; and (iii) Matplotlib for plotting and visualisation and as an alternative for loading and saving images.

Note: In contrast to the Anaconda distribution, Miniconda supplies only the package management system. This allows to set up minimalistic environments when the disk space is limited (e.g., pool PCs offered by ISP). If you have Anaconda already installed and would like to use it instead, the instructions below should also work.

Setup Miniconda:

- Install Miniconda following the instructions at <https://docs.conda.io/projects/conda/en/latest/user-guide/install/index.html>.
- To test if the installation was successful, the command `conda list` should return the list of installed packages.
- We will now create a new conda environment with the name `cv1`. We assume that the file `requirements.txt` with the dependencies we provide is in the current directory, so simply execute

```
1 conda env create -f requirements.txt -n cv1
2 source activate cv1
```

Here, the second command activates the environment `cv1`.

Tasks: We are now ready to execute some Python code! You can use the entry-point `main.py` to import the functions from `problem1.py` to test your code. As a warm-up, complete the following tasks in `problem1.py`.

- Load image `data/a1p1.png` using the `load_image` function we provide in `main.py`.
- Implement `display_image` to show the image using `matplotlib`.
- Implement `save_as_npy` to save the image (a numpy array) to a binary `.npy` file (use `numpy.save()`).
- Implement `load_npy` to load the previously saved file. Check that it is the same image with `display_image`.
- Implement `mirror_horizontal` to horizontally mirror (flip) the image, i.e., the resulting image should revert the pixel order in the horizontal direction.
- Implement `display_images` to display the original and the mirrored (flipped) image from the previous step in *one* plot.

Submission: Please only include the file `problem1.py` in your submission.

Problem 2: Bayer Interpolation (10 Points)

Although today's digital cameras output RGB color pictures or videos, most of them do not have three separate RGB sensing chips but only one single chip based on a *color filter array* (CFA), where alternating sensor sites are covered by different color filters. The Bayer pattern (see Fig. 1) is the most commonly used pattern. In this problem the task is to construct a RGB color image from the data in the Bayer pattern through interpolating the missing color values.

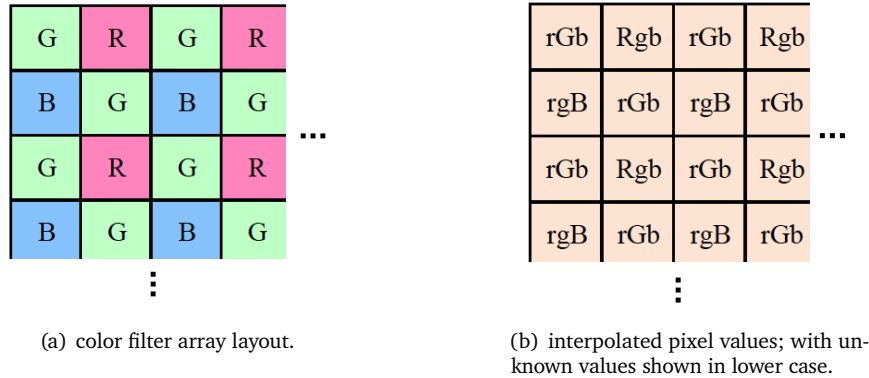


Figure 1: Bayer RGB pattern.

Your task is to write Python code that uses *bilinear interpolation* to restore all missing RGB color values (Fig. 1(b)) from the image data in the Bayer pattern (Fig. 1(a)). Load the image data saved in the Bayer pattern from `bayerdata.npy` and show two images:

1. Display the RGB color image directly transformed from the Bayer pattern, where missing color values should be filled with 0.
2. Display the interpolated full RGB color image.

Tasks:

- Implement the data loading in the function `loadbayer`.
- Separate the color channels filling missing values with zero in the function `separatechannels`.
- Implement bilinear interpolation for the missing values in the function `interpolate`. You should use `ndimage.convolve` from `scipy` package and find appropriate 3×3 kernels. For boundary pixels that might not be (bi)linearly interpolated, please take a look at the `mode` argument of `convolve` and select an appropriate option to make the color of each pixel look consistent with its neighbours.

Notes and Tips: The skeleton is given in `problem2.py`. Please, do not change the function signatures, that is the number of variables the function has as the arguments and the return values. If you are not familiar with (bilinear) interpolation, you can find a tutorial here:

<http://www.cambridgeincolour.com/tutorials/image-interpolation.htm>.

Submission: Please only include the file `problem2.py` in your submission.

Problem 3: Intrinsic Calibration (10 Points)

In this exercise, you will find the projection matrix of a camera and its intrinsic and extrinsic components – the process called *camera calibration*. We provide 2D/3D point correspondences extracted from a calibration pattern similar to the one shown in the Fig. 2. Your task is to use the *homogeneous least squares* approach discussed in the lecture to recover the projection parameters.




Figure 2: Example of a calibration target

Tasks:

- Implement the function `load_points` to load the 2D and 3D point correspondences from a numpy file `points.npz`. N correspondences are stored as $N \times 3$ and $N \times 4$ arrays and can be accessed via `image` and `world` keys.
- Construct matrix A , as defined in the lecture slides, in the function `create_A`. Recall that this matrix emerges from the cross-product between *homogeneous* representation of the 2D points and their projected 3D correspondences.
- We will now use homogeneous least squares to find the elements in our projections matrix P in function `homogeneous_Ax`. Perform Singular Value Decomposition (SVD) of P and compute the right singular vector corresponding to the smallest singular value in function. Function `homogeneous_Ax` should return the reshaped vector into 3×4 matrix, our projection matrix P .
- Now that we have obtained matrix P , we would like to find the intrinsics K , as well as the rotation R and translation t of the camera frame to the world frame. Implement function `solve_KR` that uses RQ-decomposition of matrix P and returns matrices K and R . Note that RQ-decomposition is not unique, and you might need to flip the signs in K to ensure the elements (focal length) are positive. In this case, remember to flip the sign in matrix R accordingly.
- To find the translation component t of the extrinsics, we will first find the shift c between the coordinate frames of the camera and the world, as we defined it in the lecture. Implement `solve_c` that uses SVD of P to compute its nullspace and returns c in a *non-homogeneous* form (3D vector).

Submission: Please only include your writeup of `problem3.py` in your submission.

Problem 4: Image Filtering and Edge Detection (15 Points)

In this problem you will get to know how to compute image derivatives and create a simple edge detector. For testing use image `a1p4.png`.

Tasks:

- Implement `createfilters` that creates a 3×3 x -derivative filter that computes the derivative using central differences (i.e., $\frac{f(x+h;y)-f(x-h;y)}{2h}$ where $h = 1$) and at the same time smoothes in y -direction with a 3×1 Gaussian filter of $\sigma = 0.9$. In other terms, you obtain the full 3×3 filter by combining a 1×3 with a 3×1 filter. Also, create a corresponding y -derivative filter. Make sure that the filters are correctly normalized.
- Make use of the created filters in `filterimage` to compute the x -derivatives ∂x and y -derivatives ∂y of the input image from above. You can use the `scipy` method `ndimage.convolve` to do the filtering. Note: Make sure that the sign of the derivatives is correct.
- Compute and display the gradient magnitude $\sqrt{\partial_x^2 + \partial_y^2}$. Detect edges by thresholding the gradient magnitude in `detectededges`, i.e., check if the gradient magnitude at each pixel exceeds a certain value.
- Implement *non-maximum suppression* in `nonmaxsupp` to thin the obtained edges. For every edge candidate pixel, consider the grid direction (out of 4 directions) that is “most orthogonal” to the edge. If one of the two neighbors in this direction has a larger gradient magnitude, remove the central pixel from the edge map.

The initial threshold is set to 1.0, however, this is quite an arbitrary choice. Experiment with various threshold values and choose one that shows the “important” image edges. Briefly explain with a comment in the code how you found this threshold and why you chose it.

Submission: Please only include the file `problem4.py` in your submission.